6 Big Data File Formats Compared, Pt. 3

§

🔷 Made with Obsidian

📇 Type blog   ⬡ Category big-data   </> Technologies Python   📖 Website Post Link

Over the last two articles of this series, we have discussed different Big Data file formats and their overall characteristics. We have also performed writing & reading examples using different Python modules & methods.

In this section, we will focus on comparing the performance of the formats reviewed. We will evaluate writing times, reading times, and file sizes.

We will then conclude this series by reviewing specific use cases for each one of the formats, as well as discussing some recommendations.

This section will be longer than the previous ones and involve more code. We'll be using Python scripts which can be found in the Blog Article Repo.

§

Table of Contents

§

Importing the required modules

For this section, we will be using the following Python libraries and modules:

CODE

```python
# File manipulation modules
import pandas as pd
from fastavro import reader, writer, parse_schema
import pickle
import openpyxl

# Performance measurement modules
import time

# System utility modules
import os
import shutil
from pathlib import Path

# Plotting modules
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
```

§

Defining plot parameters

Since we'll be plotting our experimental results, we will need to define our plot parameters beforehand:

CODE

```
# Before anything else, delete the Matplotlib
# font cache directory if it exists, to ensure
# custom font proper loading
try:
    shutil.rmtree(matplotlib.get_cachedir())
except FileNotFoundError:
    pass

# Define main color as hex
color_main = '#1a1a1a'

# Define title & label padding
text_padding = 18

# Define font sizes
title_font_size = 17
label_font_size = 14

# Define rc params
plt.rcParams['figure.figsize'] = [14.0, 7.0]
plt.rcParams['figure.dpi'] = 300
plt.rcParams['grid.color'] = 'k'
plt.rcParams['grid.linestyle'] = ':'
plt.rcParams['grid.linewidth'] = 0.5
plt.rcParams['font.family'] = 'sans-serif'
plt.rcParams['font.sans-serif'] = ['Lora']
```

--------------------------------- § ---------------------------------

Preparing the data set

To reduce noise in our performance measurements, we will employ a slightly larger file for this section. We will be using the `Spotify Charts` data set (`2.48 GB`) published by `DHRUVIL DAVE` , which you can download from Kaggle.

We can start by reading the data set as a `pandas.DataFrame` object:

CODE

```
# Load csv file as pandas.DataFrame object
df = pd.read_csv('datasets/charts.csv')
```

This process should take a couple of minutes, depending on the specifications of each machine.

Upon concluding, we should end up with a `pandas.DataFrame` object `df` with the following shape:

CODE

```
# Get the shape of the object
df.shape
```

```
(26173514, 9)
```

Meaning *26,173,514* rows by *9* columns.

Next, we will need to do some preprocessing before beginning with the tests:

CODE

```python
# Singe we have nan values, we will remove them
df = df.dropna()

# Check the current data types and see if casting is required
df.dtypes

# Cast to required data types
# Date is currently a string. We will cast it to Pandas DateTime in integer format.
# since Avro does not support original DateTime
df['date'] = pd.to_datetime(df['date'])
df['date'] = df['date'].apply(lambda x: x.value)

# Streams will be casted to integer type
df['streams'] = df['streams'].astype('int')

# Finally, reset index since .feather does not support
# serializing a non-default index
df = df.reset_index(drop = True)
```

We converted our `date` field to an `int` `datetime64` data type because the `fastavro` library does not currently support `str` `datetime64` objects.

Since we will be defining a schema for this file format, a conventional `datetime64` data type would raise a `TypeError`.

Lastly, we will create two new directories:

- The first one for storing all of our written files.
- The second one for storing all of our performance test results.

CODE

```
mkdir performance_tests
mkdir performance_results
```

Now, we're ready to start designing our experiment.

§

Experiment design

For both writing & reading performance tests, we will be using a collection of measurements per file format. This is always a good practice in experimental design and will help us calculate a complete set of descriptive statistics.

Our experiment for the **writing process** will consist of the following steps:

1. Define the variables to measure.
2. Set the number of trials as a control variable.
3. Begin with the first trial, measuring variables of interest.
4. Store measurements.
5. Repeat for the other file formats.
6. Consolidate results and perform a statistical description.
7. Delete all written files except one to read in the next experiment.

Our experiment for the **reading process** will consist of the following steps:

1. Define the variables to measure.
2. Set the number of trials as a control variable.
3. Begin with the first trial, measuring variables of interest.
4. Store measurements.
5. Clear memory
6. Repeat for the other file formats.
7. Consolidate results and perform a statistical description.

—————————————— § ——————————————

Performance tests

1. Parameter definition

We will start by defining our experiment parameters:

CODE

```
# Define number of trials n
n = 20

# Define performance tests output path
path = 'performance_tests/'
```

2. Writing performance tests

We will start by defining our `writingPerformance()` function, which will accept the following parameters :

- `n` : `int`
  - Number of trials.
- `path` : `str`
  - Path for writing results.
- `df` : `pandas.DataFrame`
  - DataFrame Object containing the data set.

Upon calling, it will return the following:

- `measured_vars_w` : `dict`
  - Execution time for each file format with `n` number of trials.

Once we have the expected inputs and outputs, the idea is to perform the following:

1. Declare a `pandas.core.series.Series` for each file format, where we will store writing times. This will result in 8 objects in total.
2. Declare an empty dictionary `measured_vars_w`, used to store key-value pairs of file format-Pandas series of measurements.
3. Define a `for` loop to iterate over the number of trials `n`.
4. Define a writing method for each file format.
5. Set a `start` timer variable before writing execution using the `time.time()` method. This will be our initial time marker.
6. Set an `end` timer variable after writing execution using the `time.time()` method. This will be our final time marker.
7. Calculate the `exec_time` for each loop by calculating the difference between `end` and `start` variables.
8. Append the `exec_time` to our corresponding `pandas.core.series.Series` object.
9. Upon loop completion, assign a key-value pair to our `measured_vars_w` dictionary corresponding to the file format and measured execution times.
10. Remove all generated files except one using the `os.remove` method.
11. Repeat for all remaining file formats.

We can translate our pseudocode into code:

CODE

```python
def writingPerformance(n, path, df):
    '''
    Parameters
    ----------
    n : int
        Number of trials.
    path : str
        Path for result writing.

    df : pandas.DataFrame
        DataFrame Object containing data set.

    Returns
    -------
    measured_vars_w : dict
        Execution time for each file format, with n number of trials.

    '''

    # Declare pd.Series() for storing the measured variables for each file format
    wtime_csv = pd.Series([], dtype='float64')
    wtime_txt = pd.Series([], dtype='float64')
    wtime_feather = pd.Series([], dtype='float64')
    wtime_parquet_NP = pd.Series([], dtype='float64')
    wtime_parquet_SP = pd.Series([], dtype='float64')
    wtime_parquet_MP = pd.Series([], dtype='float64')
    wtime_avro = pd.Series([], dtype='float64')
    wtime_pickle = pd.Series([], dtype='float64')

    # Declare a dictionary for storing all series
    measured_vars_w = {}

    # -------------------------------
    # 1. CSV
    # -------------------------------
    for trial in range(n):

        # Start trial
        print(f'CSV trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Write to file
        df.to_csv(path + 'CSV' + '_' + str(trial) + '.csv')

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        wtime_csv = pd.concat([pd.Series([exec_time]), wtime_csv])

    # Define series title
```

```python
    wtime_csv.name = 'Writing Time [s]'

    # Add measurements to dictionary
    measured_vars_w['01_CSV'] = wtime_csv.reset_index(drop = True)

    # Remove all files from dir except one
    for trial in range(n - 1):
        os.remove(path + 'CSV' + '_' + str(trial) + '.csv')

    # ------------------------------
    # 2. TXT
    # ------------------------------
    for trial in range(n):

        # Start trial
        print(f'TXT trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Write to file
        df.to_csv(path + 'TXT' + '_' + str(trial) + '.txt', sep = '\t')

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        wtime_txt = pd.concat([pd.Series([exec_time]), wtime_txt])

    # Define series title
    wtime_txt.name = 'Writing Time [s]'

    # Add measurements to dictionary
    measured_vars_w['02_TXT'] = wtime_txt.reset_index(drop = True)

    # Remove all files from dir except one
    for trial in range(n - 1):
        os.remove(path + 'TXT' + '_' + str(trial) + '.txt')

    # ------------------------------
    # 3. Feather
    # ------------------------------
    for trial in range(n):

        # Start trial
        print(f'Feather trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Write to file
        df.to_feather(path + 'Feather' + '_' + str(trial) + '.feather')
```

```python
        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        wtime_feather = pd.concat([pd.Series([exec_time]), wtime_feather])

    # Define series title
    wtime_feather.name = 'Writing Time [s]'

    # Add measurements to dictionary
    measured_vars_w['03_Feather'] = wtime_feather.reset_index(drop = True)

    # Remove all files from dir except one
    for trial in range(n - 1):
        os.remove(path + 'Feather' + '_' + str(trial) + '.feather')

    # -------------------------------
    # 4. Parquet non-partitioned
    # -------------------------------
    for trial in range(n):

        # Start trial
        print(f'Parquet NP trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Write to file
        df.to_parquet(path + 'Parquet_NP' + '_' + str(trial) + '.parquet')

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        wtime_parquet_NP = pd.concat([pd.Series([exec_time]), wtime_parquet_NP])

    # Define series title
    wtime_parquet_NP.name = 'Writing Time [s]'

    # Add measurements to dictionary
    measured_vars_w['04_Parquet_NP'] = wtime_parquet_NP.reset_index(drop = True)

    # Remove all files from dir except one
    for trial in range(n - 1):
        os.remove(path + 'Parquet_NP' + '_' + str(trial) + '.parquet')

    # -------------------------------
    # 5. Parquet single-partitioned
    # -------------------------------
    for trial in range(n):
```

```python
        # Start trial
        print(f'Parquet SP trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Write to file
        df.to_parquet(path + 'Parquet_SP' + '_' + str(trial) + '.parquet', partition_cols =
['region'])

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        wtime_parquet_SP = pd.concat([pd.Series([exec_time]), wtime_parquet_SP])

    # Define series title
    wtime_parquet_SP.name = 'Writing Time [s]'

    # Add measurements to dictionary
    measured_vars_w['05_Parquet_SP'] = wtime_parquet_SP.reset_index(drop = True)

    # Remove all files from dir except one
    for trial in range(n - 1):
        shutil.rmtree(path + 'Parquet_SP' + '_' + str(trial) + '.parquet')

    # ------------------------------
    # 6. Parquet multi-partitioned
    # ------------------------------
    for trial in range(n):

        # Start trial
        print(f'Parquet MP trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Write to file
        df.to_parquet(path + 'Parquet_MP' + '_' + str(trial) + '.parquet', partition_cols =
['region', 'trend'])

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        wtime_parquet_MP = pd.concat([pd.Series([exec_time]), wtime_parquet_MP])

    # Define series title
    wtime_parquet_MP.name = 'Writing Time [s]'
```

```python
    # Add measurements to dictionary
    measured_vars_w['06_Parquet_MP'] = wtime_parquet_MP.reset_index(drop = True)

    # Remove all files from dir except one
    for trial in range(n - 1):
        shutil.rmtree(path + 'Parquet_MP' + '_' + str(trial) + '.parquet')


    # -------------------------------
    # 7. Avro
    # -------------------------------
    # Define the schema
    schema = {
        'type': 'record',
        'name': 'performance_comp',
        'namespace': 'performance_comp',
        'doc': 'This schema consists of 2 int types, 1 datetime type and 6 string types',
        'fields': [
            {'name': 'title', 'type': 'string'},
            {'name': 'rank', 'type': 'int'},
            {'name': 'date', 'type': 'long'},
            {'name': 'artist', 'type': 'string'},
            {'name': 'url', 'type': 'string'},
            {'name': 'region', 'type': 'string'},
            {'name': 'chart', 'type': 'string'},
            {'name': 'trend', 'type': 'string'},
            {'name': 'streams', 'type': 'int'}
        ]
    }

    # Parse the schema
    parsed_schema = parse_schema(schema)

    # Convert pd.DataFrame to records (list of dictionaries)
    records = df.to_dict('records')

    for trial in range(n):

        # Start trial
        print(f'Avro trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Write to Avro file
        with open(path + 'Avro' + '_' + str(trial) + '.avro', 'wb') as file:
            writer(file, parsed_schema, records)

        file.close()

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start
```

```python
        # Append time to series
        wtime_avro = pd.concat([pd.Series([exec_time]), wtime_avro])

    # Define series title
    wtime_avro.name = 'Writing Time [s]'

    # Add measurements to dictionary
    measured_vars_w['07_Avro'] = wtime_avro.reset_index(drop = True)

    # Remove all files from dir except one
    for trial in range(n - 1):
        os.remove(path + 'Avro' + '_' + str(trial) + '.avro')


    # -----------------------------
    # 8. Pickle open file
    # -----------------------------
    for trial in range(n):

        # Start trial
        print(f'Pickle trial {trial} of {n} started...')

        # Convert pd.DataFrame to records (list of dictionaries)
        records = df.to_dict('records')

        # Start timer
        start = time.time()

        file = open(path + 'Pickle' + '_' + str(trial) + '.pickle', 'wb')

        # Write open file to disk
        pickle.dump(records, file)

        file.close()

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        wtime_pickle = pd.concat([pd.Series([exec_time]), wtime_pickle])

    # Define series title
    wtime_pickle.name = 'Writing Time [s]'

    # Add measurements to dictionary
    measured_vars_w['08_Pickle'] = wtime_pickle.reset_index(drop = True)

    # Remove all files from dir except one
    for trial in range(n - 1):
        os.remove(path + 'Pickle' + '_' + str(trial) + '.pickle')

    return measured_vars_w
```

3. Reading performance tests

Similar to the writing performance tests section, we will start by defining our `readingPerformance()` function, which will accept the following parameters :

- `n` : `int`
  - Number of trials.
- `path` : `str`
  - Path for writing results.
- `df` : `pandas.DataFrame`
  - DataFrame Object containing data set.

Upon calling, it will return the following:

- `measured_vars_r` : `dict`
  - Execution time for each file format, with `n` number of trials.

Once we have the expected inputs and outputs, the idea is to perform the following:

1. Declare a `pandas.core.series.Series` for each file format, where we will store writing times. This would result in 8 objects in total.
2. Declare an empty dictionary `measured_vars_w`, used to store key-value pairs of file format-Pandas series of measurements.
3. Define a `for` loop to iterate over the number of trials `n`.
4. Define a reading method for each file format.
5. Set a `start` timer variable before writing execution using the `time.time()` method. This will be our initial time marker.
6. Set an `end` timer variable after writing execution using the `time.time()` method. This will be our final time marker.
7. Calculate the `exec_time` for each loop by calculating the difference between `end` and `start` variables.
8. Append the `exec_time` to our corresponding `pandas.core.series.Series` object.
9. Delete the read object from memory.
10. Upon loop completion, assign a key-value pair to our `measured_vars_w` dictionary corresponding to the file format and measured execution times.
11. Repeat for all remaining file formats.

We can translate our pseudocode into code:

CODE

```python
def readingPerformance(n, path, df):
    '''

    Parameters
    ----------
    n : int
        Number of trials.
    path : str
        Path for result writing.

    df : pandas.DataFrame
        DataFrame Objectn containing data set.

    Returns
    -------
    measured_vars_r : dict
        Execution time for each file format, with n number of trials.

    '''

    # Declare pd.Series() for storing the measured variables for each file format
    rtime_csv = pd.Series([], dtype='float64')
    rtime_txt = pd.Series([], dtype='float64')
    rtime_feather = pd.Series([], dtype='float64')
    rtime_parquet_NP = pd.Series([], dtype='float64')
    rtime_parquet_SP = pd.Series([], dtype='float64')
    rtime_parquet_MP = pd.Series([], dtype='float64')
    rtime_avro = pd.Series([], dtype='float64')
    rtime_pickle = pd.Series([], dtype='float64')

    # Declare a dictionary for storing all series
    measured_vars_r = {}

    # -------------------------------
    # 1. CSV
    # -------------------------------
    for trial in range(n):

        # Start trial
        print(f'CSV trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Read from file
        df = pd.read_csv(path + 'CSV' + '_' + str(n-1) + '.csv')

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        rtime_csv = pd.concat([pd.Series([exec_time]), rtime_csv])

        # Delete file from memory
```

```python
        del df

    # Define series title
    rtime_csv.name = 'Reading Time [s]'

    # Add measurements to dictionary
    measured_vars_r['01_CSV'] = rtime_csv.reset_index(drop = True)

    # -------------------------------
    # 2. TXT
    # -------------------------------
    for trial in range(n):

        # Start trial
        print(f'TXT trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Read from file
        df = pd.read_csv(path + 'TXT' + '_' + str(n-1) + '.txt', sep = '\t')

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        rtime_txt = pd.concat([pd.Series([exec_time]), rtime_txt])

        # Delete file from memory
        del df

    # Define series title
    rtime_txt.name = 'Reading Time [s]'

    # Add measurements to dictionary
    measured_vars_r['02_TXT'] = rtime_txt.reset_index(drop = True)

    # -------------------------------
    # 3. Feather
    # -------------------------------
    for trial in range(n):

        # Start trial
        print(f'Feather trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Read from file
        df = pd.read_feather(path + 'Feather' + '_' + str(n-1) + '.feather')

        # End timer
        end = time.time()
```

```python
    # Calculate execution time
    exec_time = end - start

    # Append time to series
    rtime_feather = pd.concat([pd.Series([exec_time]), rtime_feather])

    # Delete file from memory
    del df

# Define series title
rtime_feather.name = 'Reading Time [s]'

# Add measurements to dictionary
measured_vars_r['03_Feather'] = rtime_feather.reset_index(drop = True)


# ------------------------------
# 4. Parquet non-partitioned
# ------------------------------
for trial in range(n):

    # Start trial
    print(f'Parquet NP trial {trial} of {n} started...')

    # Start timer
    start = time.time()

    # Read from file
    df = pd.read_parquet(path + 'Parquet_NP' + '_' + str(n-1) + '.parquet')

    # End timer
    end = time.time()

    # Calculate execution time
    exec_time = end - start

    # Append time to series
    rtime_parquet_NP = pd.concat([pd.Series([exec_time]), rtime_parquet_NP])

    # Delete file from memory
    del df

# Define series title
rtime_parquet_NP.name = 'Reading Time [s]'

# Add measurements to dictionary
measured_vars_r['04_Parquet_NP'] = rtime_parquet_NP.reset_index(drop = True)


# ------------------------------
# 5. Parquet single-partitioned
# ------------------------------
for trial in range(n):

    # Start trial
    print(f'Parquet SP trial {trial} of {n} started...')
```

```python
        # Start timer
        start = time.time()

        # Read from file
        df = pd.read_parquet(path + 'Parquet_SP' + '_' + str(n-1) + '.parquet')

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        rtime_parquet_SP = pd.concat([pd.Series([exec_time]), rtime_parquet_SP])

        # Delete file from memory
        del df

    # Define series title
    rtime_parquet_SP.name = 'Reading Time [s]'

    # Add measurements to dictionary
    measured_vars_r['05_Parquet_SP'] = rtime_parquet_SP.reset_index(drop = True)


    # -------------------------------
    # 6. Parquet multi-partitioned
    # -------------------------------
    for trial in range(n):

        # Start trial
        print(f'Parquet MP trial {trial} of {n} started...')

        # Start timer
        start = time.time()

        # Read from file
        df = pd.read_parquet(path + 'Parquet_MP' + '_' + str(n-1) + '.parquet')

        # End timer
        end = time.time()

        # Calculate execution time
        exec_time = end - start

        # Append time to series
        rtime_parquet_MP = pd.concat([pd.Series([exec_time]), rtime_parquet_MP])

        # Delete file from memory
        del df

    # Define series title
    rtime_parquet_MP.name = 'Reading Time [s]'

    # Add measurements to dictionary
    measured_vars_r['06_Parquet_MP'] = rtime_parquet_MP.reset_index(drop = True)
```

```python
# -------------------------------
# 7. Avro
# -------------------------------

for trial in range(n):

    # Start trial
    print(f'Avro trial {trial} of {n} started...')

    # Define list of dictionaries
    lod = []

    # Start timer
    start = time.time()

    # Read from Avro file
    with open(path + 'Avro' + '_' + str(n-1) + '.avro', 'rb') as fo:
        avro_reader = reader(fo)

        for record in avro_reader:
            lod.append(record)

    # Close the BufferedReader object
    fo.close()

    # End timer
    end = time.time()

    # Calculate execution time
    exec_time = end - start

    # Append time to series
    rtime_avro = pd.concat([pd.Series([exec_time]), rtime_avro])

    # Delete file from memory
    del lod

# Define series title
rtime_avro.name = 'Reading Time [s]'

# Add measurements to dictionary
measured_vars_r['07_Avro'] = rtime_avro.reset_index(drop = True)

# -------------------------------
# 8. Pickle open file
# -------------------------------
for trial in range(n):

    # Start trial
    print(f'Pickle trial {trial} of {n} started...')

    # Start timer
    start = time.time()

    # Read from file
    with open(path + 'Pickle' + '_' + str(n-1) + '.pickle', 'rb') as file:
```

```python
            my_pickled_object = pickle.load(file)

            # Close the BufferedReader object
            file.close()

            # End timer
            end = time.time()

            # Calculate execution time
            exec_time = end - start

            # Append time to series
            rtime_pickle = pd.concat([pd.Series([exec_time]), rtime_pickle])

            # Delete file from memory
            del my_pickled_object

        # Define series title
        rtime_pickle.name = 'Reading Time [s]'

        # Add measurements to dictionary
        measured_vars_r['08_Pickle'] = rtime_pickle.reset_index(drop = True)

        return measured_vars_r
```

4. Analysis

Once we have both the `writingPerformance()` and `readingPerformance()` functions declared, we can define an `analysis()` function which will accept the following parameters:

- `n` : `int`
  - Number of trials.
- `path` : `str`
  - Path for writing results.
- `df` : `pandas.DataFrame`
  - DataFrame Object containing data set.

Upon calling, it will return the following:

- `measured_vars_w` : `dict`
  - Writing time of each file format with `n` number of trials.
- `measured_vars_r` : `dict`
  - Reading time of each file format with `n` number of trials.
- `stat_dw` : `dict`
  - Statistical description of measured writing times for all formats.
- `stat_dr` : `dict`
  - Statistical description of measured reading times for all formats.
- `size_d` : `dict`
  - Statistical description of measured file/folder sizes for all formats.

Once we have the expected inputs and outputs, the idea is to perform the following:

4.1 Writing performance analysis

1. Define a list `tests` containing the names of each file format. We will use this object as iterable.
2. Call the `writingPerformance()` function and assign it to a `measured_vars_w` object.
3. Define an empty dictionary `stat_dw` for saving the statistical descriptions for the writing test.
4. Iterate over the `tests` object, indexing the `measured_vars_w` dictionary on each loop and calculating its statistical description.
5. Assign a key-value pair on each loop, consisting of the file format name as the key and the statistical description object as the value.

4.2 Reading performance analysis

1. Call the `readingPerformance()` function and assign it to a `measured_vars_r` object.
2. Define an empty dictionary `stat_dr` for saving the statistical descriptions for the reading test.
3. Iterate over the `tests` object, indexing the `measured_vars_r` dictionary on each loop and calculating its statistical description.
4. Assign a key-value pair on each loop, consisting of the file format name as the key and the statistical description object as the value.

4.3 File size analysis

1. Declare an empty dictionary `size_d` for storing the file size values.
2. For each file, calculate its size using the `os.path.getsize` method.
3. For each folder (parquet partitions), recursively calculate its size using the `f.stat().st_size` method.
4. Convert file sizes:
    1. The `os.path.getsize` method returns the file size in bytes, so we need to divide the result by $1,024^2$ to get our measurements in `MB` units.
    2. The `f.stat().st_size` method returns the folder size in bytes, so we need to divide the result by $1,024^2$ to get our measurements in `MB` units.

We can translate our pseudocode into code:

CODE

```python
def analysis(n, path, df):
    '''
    Parameters
    ----------
    n : int
        Number of trials.
    path : str
        Path for result writing.

    df : pandas.DataFrame
        DataFrame Objectn containing data set.

    Returns
    -------
    measured_vars_w : dict
        Writing time of each file format, with n number of trials.

    measured_vars_r : dict
        Reading time of each file format, with n number of trials.

    stat_dw : dict
        Statistical description of measured writing times for all formats.

    stat_dr : dict
        Statistical description of measured reading times for all formats.

    size_d : dict
        Statistical description of measured file/folder sizes for all formats.
    '''

    # Define all tests for including in statistical description
    tests = ['01_CSV',
             '02_TXT',
             '03_Feather',
             '04_Parquet_NP',
             '05_Parquet_SP',
             '06_Parquet_MP',
             '07_Avro',
             '08_Pickle']


    # -------------------------------
    # Writing Analysis
    # -------------------------------

    # Perform writing experiment and get measured_vars_w dictionary
    measured_vars_w = writingPerformance(n, path, df)

    # Define statistical results dict
    stat_dw = {}

    # Statistical Description for Writing

    # Extract each set of values, describe them and save them in a new dict
    for i in tests:
        stat_v = measured_vars_w[i].describe()
        stat_dw[i] = stat_v
```

```python
    # ------------------------------
    # Reading Analysis
    # ------------------------------

    # Perform reading experiment and get measured_vars_r dictionary
    measured_vars_r = readingPerformance(n, path, df)

    # Define statistical results dict
    stat_dr = {}

    # Statistical Description for Reading

    # Extract each set of values, describe them and save them in a new dict
    for i in tests:
        stat_v = measured_vars_r[i].describe()
        stat_dr[i] = stat_v

    # ------------------------------
    # File Size Analysis
    # ------------------------------

    # Define size results dict
    size_d = {}

    # Calculate file & dir sizes
    size_d['01_CSV'] = os.path.getsize(path + 'CSV' + '_' + str(n-1) + '.csv') / (1024**2)
    size_d['02_TXT'] = os.path.getsize(path + 'TXT' + '_' + str(n-1) + '.txt') / (1024**2)
    size_d['03_Feather'] = os.path.getsize(path + 'Feather' + '_' + str(n-1) + '.feather') /
(1024**2)
    size_d['04_Parquet_NP'] = os.path.getsize(path + 'Parquet_NP' + '_' + str(n-1) + '.parquet')
/ (1024**2)
    path_Parquet_SP = Path(path + 'Parquet_SP_19.parquet')
    size_d['05_Parquet_SP'] = sum(f.stat().st_size for f in path_Parquet_SP.glob('**/*') if
f.is_file()) / (1024**2)
    path_Parquet_MP = Path(path + 'Parquet_MP_19.parquet')
    size_d['06_Parquet_MP'] = sum(f.stat().st_size for f in path_Parquet_MP.glob('**/*') if
f.is_file()) / (1024**2)
    size_d['07_Avro'] = os.path.getsize(path + 'Avro' + '_' + str(n-1) + '.avro') / (1024**2)
    size_d['08_Pickle'] = os.path.getsize(path + 'Pickle' + '_' + str(n-1) + '.pickle') /
(1024**2)

    # Return a dictionary including the actual measured time values of all methods
    # Return a dictionary including the statistical description of all methods
    return measured_vars_w, measured_vars_r, stat_dw, stat_dr, size_d
```

If we closely examine lines `1049` through `1059`, we see that we need to declare different methods for calculating file sizes & folder sizes.

We can then call our `analysis()` function and assign the outputs to 5 different objects:

CODE

```
# Call analysis function
measured_vars_w, measured_vars_r, stat_dw, stat_dr, size_d = analysis(n, path, df)

# Print the type and shape of each object
print(type(measured_vars_w))
print(len(measured_vars_w))

print(type(measured_vars_r))
print(len(measured_vars_r))

print(type(stat_dw))
print(len(stat_dw))

print(type(stat_dr))
print(len(stat_dr))

print(type(size_d))
print(len(size_d))
```

OUTPUT

```
<class 'dict'>
8
<class 'dict'>
8
<class 'dict'>
8
<class 'dict'>
8
<class 'dict'>
8
```

The shapes of our objects match the output requirements, as we have eight tested file formats.

§

Performance results

We now have everything we need to make sense of the data we just generated.

In this example, the quantity of data is manageable since we only have eight file formats and 20 measurements per format. Regardless, it is always a good practice to visualize the results in a plot as a first step.

1. Plotting the results

We'll proceed with some visualization techniques appropriate for the results we have.

1.1 Bar chart for file sizes

A bar chart is a simple way to visualize data quickly. We can generate a bar chart using the `matplotlib.pyplot` module:

CODE

```python
# Create figure
plt.figure('File Sizes Bar Chart')

# Plot the file sizes
plt.bar(size_d.keys(), size_d.values(), color = color_main)

# Enable grid
plt.grid(True, zorder=0)

# Set xlabel and ylabel
plt.xlabel('File Format', fontsize=label_font_size, labelpad=text_padding)
plt.ylabel('File Size [GB]', fontsize=label_font_size, labelpad=text_padding)

# Remove bottom and top separators
sns.despine(bottom=True)

# Add plot title
plt.title('File/Folder Sizes in Gigabytes', fontsize=title_font_size, pad=text_padding)

# Optional: Save the figure as a png image
plt.savefig('performance_results/' + 'file_sizes_bar_chart_tp.png', format = 'png', dpi = 300,
transparent = True)
plt.savefig('performance_results/' + 'file_sizes_bar_chart_bg.png', format = 'png', dpi = 300,
transparent = False)

# Close the figure
plt.close()
```

If we save the figure directly and then close it, the image will be written in the path we specify and will not display on our IDE.

If we examine line `1181`, we used an additional library called `seaborn` to include the additional parameter `sns.despine()`. This module is handy when dealing with statistical analysis visualizations. However, we will not use it in this particular section and will limit ourselves to using the `matplotlib.pyplot` module.
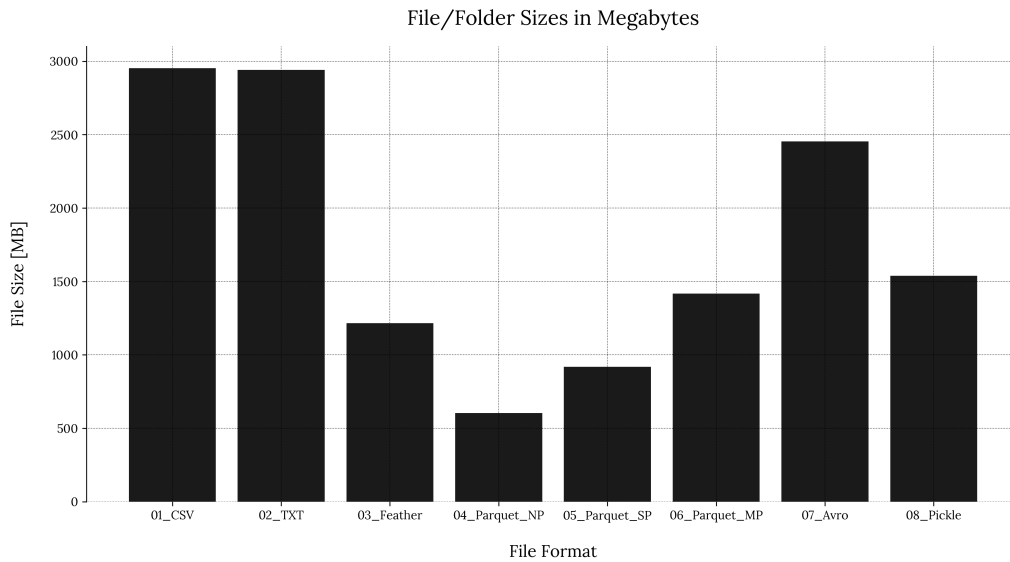
OUTPUT

File/Folder Sizes in Megabytes

FIGURE 1.1: BAR CHART DENOTING FILE/FOLDER SIZES IN `MB` *FOR EACH FILE FORMAT*

1.2 Boxplot for writing times

A boxplot is a visualization method widely used in Data Science & statistical analysis. Its purpose is to describe the distribution of experimental measurements, including useful visual information about the set.

A detailed discussion of the boxplot components is out of the scope of this article. That will be covered in another blog post.

We can generate a boxplot using the `matplotlib.pyplot` module:

CODE

```python
# Create figure
plt.figure('Writing Times Boxplot')

# Plot the writing times
plt.boxplot(measured_vars_w.values(),
            labels = measured_vars_w.keys(),
            showmeans=True)

# Enable grid
plt.grid(True, zorder=0)

# Set xlabel and ylabel
plt.xlabel("File Format", fontsize=label_font_size, labelpad=text_padding)
plt.ylabel("Writing Time [s]", fontsize=label_font_size, labelpad=text_padding)

# Remove bottom and top separators
sns.despine(bottom=True)

# Add plot title
plt.title('Writing Time in Seconds', fontsize=title_font_size, pad=text_padding)

# Optional: Save the figure as a png image
plt.savefig('performance_results/' + 'writing_time_scattered_boxplots_tp.png', format = 'png',
dpi = 300, transparent = True)
plt.savefig('performance_results/' + 'writing_time_scattered_boxplots_bg.png', format = 'png',
dpi = 300, transparent = False)

# Close the figure
plt.close()
```
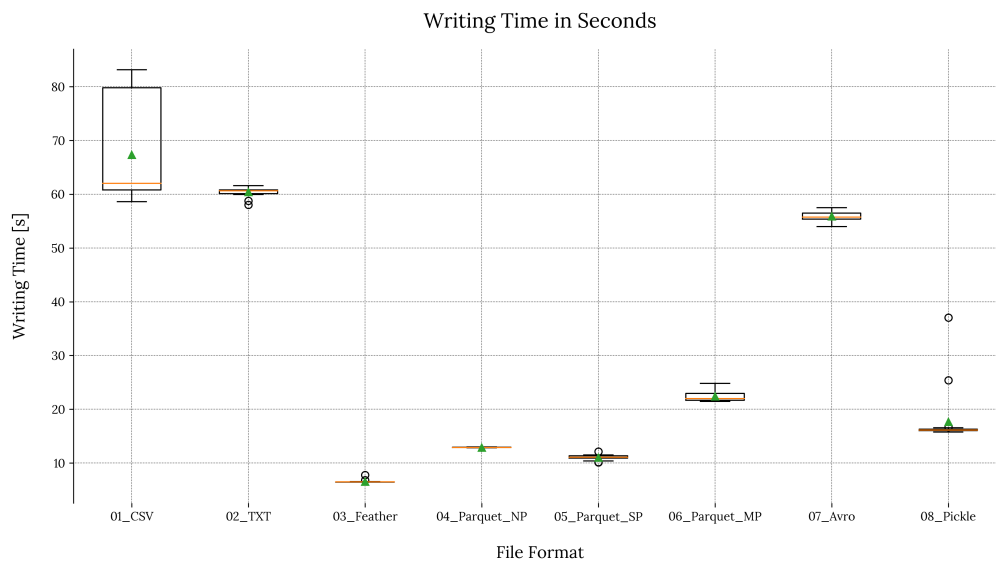
OUTPUT

FIGURE 1.2: BOXPLOT DENOTING THE DISTRIBUTION OF 20 TRIALS OF WRITING TIMES FOR EACH FILE FORMAT

1.3 Boxplot for reading times

We can perform a similar treatment to our reading time results:

CODE

```python
# Create figure
plt.figure('Reading Times Boxplot')

# Plot the writing times
plt.boxplot(measured_vars_r.values(),
            labels = measured_vars_r.keys(),
            showmeans=True)

# Enable grid
plt.grid(True, zorder=0)

# Set xlabel and ylabel
plt.xlabel("File Format", fontsize=label_font_size, labelpad=text_padding)
plt.ylabel("Reading Time [s]", fontsize=label_font_size, labelpad=text_padding)

# Remove bottom and top separators
sns.despine(bottom=True)

# Add plot title
plt.title('Reading Time in Seconds', fontsize=title_font_size, pad=text_padding)

# Optional: Save the figure as a png image
plt.savefig('performance_results/' + 'reading_time_scattered_boxplots_tp.png', format = 'png',
dpi = 300, transparent = True)
plt.savefig('performance_results/' + 'reading_time_scattered_boxplots_bg.png', format = 'png',
dpi = 300, transparent = False)

# Close the figure
plt.close()
```
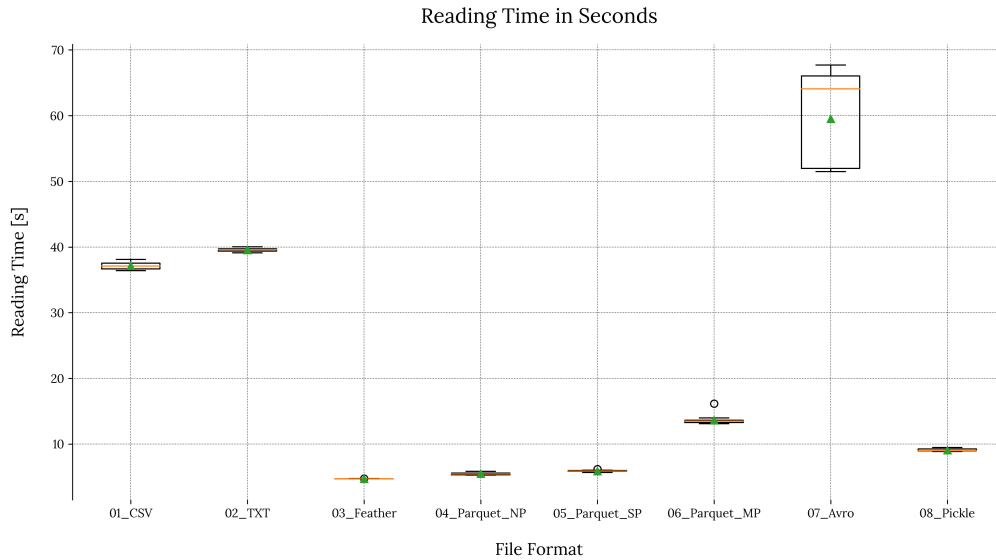
OUTPUT

Reading Time in Seconds

Figure 1.3: Boxplot Denoting The Distribution Of 20 Trials Of Reading Times For Each File Format

2. Exporting the results in a tabular format

We can also write our results in an Excel file. This is a valuable technique whenever we want to share or store information that took a fair amount of time to generate (*imagine explaining to our boss why we had to re-run a 2-hour performance test just to get the results back*).

An Excel file is also a very friendly tabular format that everyone understands. It can be used to make further analyses such as pivoting or calculating statistical measures (*mean, min, max, stdev, among others*).

For this part, we will be writing four `.xlsx` files, two for writing results and two for reading results. Each file will have eight tabs, each consisting of the file format and the writing and reading times in seconds, respectively:

CODE

```python
# Define function to export results to Excel file
def results_to_excel(dseries_dict, path):
    """Write dictionary of dataframes to separate sheets, within
        1 file."""
    writer = pd.ExcelWriter(path, engine='openpyxl')

    for tab_name, dseries in dseries_dict.items():
        dseries.to_excel(writer, sheet_name=tab_name)

    writer.close()

# Define file for writing results
path_w = 'performance_results/' + 'measured_vars_w.xlsx'
path_dw = 'performance_results/' + 'stat_dw.xlsx'

# Call function on writing results
results_to_excel(measured_vars_w, path_w)
results_to_excel(stat_dw, path_dw)

# Define file for reading results
path_r = 'performance_results/' + 'measured_vars_r.xlsx'
path_dr = 'performance_results/' + 'stat_dr.xlsx'

# Call function on reading results
results_to_excel(measured_vars_r, path_r)
results_to_excel(stat_dr, path_dr)
```

§

Side-by-side comparison

1. Consolidated results

Format	Size [MB]	Avg. Writing Time [s]	Writing Method	Avg. Reading Time [s]	Reading Method
`.csv`	2,954	67.4	`df.to_csv()`	37.1	`pd.read_csv()`
`.txt`	2,941	60.4	`df.to_csv()`	39.6	`pd.read_csv()`
`.feather`	1,216	6.6	`df.to_feather()`	4.7	`pd.read_feather()`
`.parquet`	604	12.9	`df.to_parquet()` (`NP`)	5.5	`pd.read_parquet()`
`.parquet`	919	11.0	`df.to_parquet()` (`SP`)	5.9	`pd.read_parquet()`
`.parquet`	1,417	22.4	`df.to_parquet()` (`MP`)	13.6	`pd.read_parquet()`
`.avro`	2,455	55.9	`fastavro writer()`	59.5	`fastavro reader()`
`.pickle`	1,539	17.6	`pickle.dump()`	9.1	`pickle.load()`

Note: *Keep in mind that these values vary across systems. CPU processing power, RAM capacity, and other variables directly affect reading & writing times. Please refer to the <u>Appendix</u> section for the full list of machine specifications used in this experiment.*

2. Interpretation

We can see that the `.csv` & `.txt` file formats had the largest file sizes, while the non-partitioned `.parquet` had the smallest file size. `.parquet` folder sizes increased almost linearly as we increased the number of partitions, which makes sense if we remember that a partitioned file creates a directory hierarchy beneath the main directory.

For writing times, we can see that most of our measures were more or less consistent except with the `.pickle` file format. This is denoted by the two outliers visible in *Figure 1.2* and can be due to processes starting in the background, thus reducing resources and introducing noise in the measurements. We can mitigate or at least reduce this phenomenon by tightening our experimental conditions.

Writing `.feather` files consistently consumed the least amount of time, followed by the single partitioned `.parquet` file and non-partitioned `.parquet` file consecutively. We can also clearly see that the `.csv` and `.txt` file formats were the worst performing, followed by the `.avro` format.

For reading times, the story is slightly different: Reading from the `.avro` file consumed the most amount of time (*25 seconds more than the second worst case*), along with the highest standard deviation value of all file formats (*stdev = 7, meaning 7 seconds of deviation from the mean*). This is due to the fact that the `fastavro` library provides an iterator reading object, meaning the script had to iterate over each row and append it to an object (*in our case, a list*). This additional step adds an extra layer of computational complexity, therefore increasing reading times.

In contrast, the `.feather` file format took the least amount of time to read, closely followed by the non-partitioned & single-partitioned `.parquet` files.

<div align="center">§</div>

Use cases

From the results obtained in the previous section, we can see that the `.feather` file format offers consistently fast writing & reading speeds when compared to the other file formats. Also, the file size is considerably smaller. However, as we previously mentioned, this format is not recommended for long-term storage due to its binary form instability.

A great alternative to handling big data would be the stabler `.parquet` file format: Both non-partitioned and single-partitioned forms presented low writing & reading speeds and relatively small file sizes compared to the other formats.

Even though the non-partitioned approach was the overall highest performing, it is not always the best option: when dealing with large data sets that contain multiple aggregation levels, partitioning lets us divide the aggregation levels into subfolders, making the information accessible by blocks, meaning we would not require to read the entire file to extract a single column field group. This is especially important with big data since Python reads data and stores it in memory, so a single-partitioned 1TB `.parquet` file would be practically

impossible to read (*parsing by blocks could do the trick, but then, this is what* `.parquet` *partitioning achieves more elegantly, right?*).

Another high-performing case was the `.pickle` file format, but we did not mention it as a first or even second option for two extremely relevant reasons:

- It's platform specific.
- Malicious code can be easily injected, creating potential vulnerabilities in our environment.

These two reasons alone make `.pickle` files more of a niche solution for interchanging Python objects and most definitely not apt for production environments.

Lastly, the two worst performing file formats, `.csv` and `.txt`, which ironically are the most popular, present all of the advantages that we mentioned earlier, but have substantial limitations if we want to ensure a proper and secure data writing & loading environment; the single fact that these two file formats don't preserve data types is a good-enough reason to evaluate other alternatives.

---

§

Conclusions

In summary, not all *"big data"* file formats are tailored for handling big data. This is counterintuitive, nonetheless true. Each was created with a specific purpose and is better or worse at some applications than others.

Serialized formats, for example, present several advantages over non-serialized formats. Still, the inverse also happens: serialized formats may lose stability across versions and usually consume more processing resources when serializing/deserializing objects.

Before implementing a format, especially in a production environment, it's essential to make a detailed assessment of which formats will be used and how the encoding will be handled. Also, when working with other collaborators, setting a strict standard for data formatting & handling is of vital importance and, more often than not, overlooked in favor of implementing the *"easier way"*.

---

§

Appendix

1. Experimental conditions

Below you can find a list of the parameters that were used to obtain these results:

- **Python Build:** Python 3.11.1 (tags/v3.11.1:a7a450f, Dec 6 2022, 19:58:39) MSC v.1934 64 bit (AMD64) on win32
- **Operating System:** Windows 11 Home, Version 22H2
- **Processor:** 12th Gen Intel Core i9-12900H, 2.50 GHz
- **RAM:** 64 GB DDR5

---

§

References

- [Apache Spark, Parquet Files](#)
- [Vertica, Using Partition Columns](#)
- [This Pointer, Deleting Directories Recursively](#)
- [XlsxWriter, Creating Excel files with Python and XlsxWriter](#)

--- § ---

Copyright